

# StockDataProcessor: Advanced Stock Data Processing Library

## GITHUB :

[https://github.com/quantcommunitynitrkl/data\\_preprocessing\\_OHLCV.git](https://github.com/quantcommunitynitrkl/data_preprocessing_OHLCV.git)

## Introduction

**StockDataProcessor** is a comprehensive Python library designed for financial data analysts, quantitative traders, and data scientists working with stock market time-series data. Built on top of robust libraries like `pandas`, `yfinance`, `scikit-learn`, and `matplotlib`, it provides a unified interface for downloading historical stock data, handling missing values and outliers, performing advanced imputations, and generating insightful visualizations.

This library addresses common pain points in stock data workflows:

- **Data Acquisition:** Seamless download from Yahoo Finance.
- **Data Cleaning:** Intelligent filling of missing dates and NaN values using statistical, machine learning, and smoothing techniques.
- **Outlier Management:** Multi-method detection and treatment to ensure data integrity.
- **Visualization:** Flexible plotting for exploratory analysis, including candlestick charts and interactive Plotly figures.

Whether you're backtesting trading strategies, building predictive models, or conducting market research, **StockDataProcessor** streamlines your pipeline into a modular, extensible class.

## Key Features

- **Modular Design:** Static methods for standalone use or instance-based workflows.
- **Advanced Imputation:** Supports KNN, Markov chains, Kalman filters, and more.
- **Outlier Handling:** Seven detection methods with customizable combination strategies.
- **Rich Visualizations:** 15+ graph types, from basic histograms to choropleth maps.
- **Test-Friendly:** Built-in sample dataset for quick prototyping.
- **Error-Resilient:** Extensive input validation and graceful fallbacks.

## Installation

### Prerequisites

- Python 3.8 or higher.
- Access to the internet for data downloads (via `yfinance`).

### Dependencies

Install via pip:

```
pip install yfinance pandas numpy scikit-learn matplotlib seaborn mplfinance plotly scipy pykalman
```

For development (optional):

```
pip install -r requirements-dev.txt # Includes pytest, black, etc.
```

**Note:** `pykalman` requires NumPy; ensure it's installed for Kalman filter features.

## Quick Start

```
import pandas as pd
from stockdataprocessor import StockDataProcessor # Assuming the class is in a module named stockdataprocessor

# Download data
df = StockDataProcessor.download_stock_data("AAPL", period="6mo")

# Fill missing dates
df = StockDataProcessor.fill_missing_dates(df)

# Advanced NaN filling (example config)
col_tech_map = {"close": [{"ffill", {}}, {"sma", {"window": 5}}]}
df = StockDataProcessor.fill_nan_advanced(df, col_tech_map)

# Plot a candlestick chart
StockDataProcessor.plot_graph(df, columns=["open", "high", "low", "close"], graph_type="candlestick")

# Detect and treat outliers
outliers = StockDataProcessor.detect_outliers_advanced(df, numeric_cols=["close", "volume"])
treatment_map = {"close": [{"median_replace", {}]}]}
df_treated = StockDataProcessor.treat_outliers(df, outliers, treatment_map)
```

## API Reference

The core of the library is the `StockDataProcessor` class. All methods are static for convenience, allowing use without instantiation. Below is a detailed breakdown of each method, including parameters, returns, examples, and edge cases.

### `__init__(self, df=None)`

Initializes the processor with an optional DataFrame. Useful for instance-based workflows (e.g., chaining operations).

- **Parameters:**
  - `df` (`pd.DataFrame`, optional): Input DataFrame to process.
- **Returns:** None.
- **Example:**

```
processor = StockDataProcessor(df)
# Access via instance if needed, though static methods are preferred.
```

### `download_stock_data(ticker, period="1y", interval="1d")` (Static)

Downloads historical OHLCV data from Yahoo Finance.

#### Parameters

- `ticker` (str): Stock symbol (e.g., "AAPL").
- `period` (str): Time span (e.g., "1y", "max"). See [yfinance docs](#) for options.
- `interval` (str): Data frequency (e.g., "1d", "1h").

#### Returns

- `pd.DataFrame`: Columns: `date`, `open`, `high`, `low`, `close`, `volume`.

#### Raises

- `ValueError`: Invalid ticker or download failure.

#### Edge Cases

- Multi-index columns are flattened.
- Empty results return an empty DataFrame.

#### Example

```
df = StockDataProcessor.download_stock_data("TSLA", period="3mo", interval="1d")
print(df.head()) # Displays recent data
```

`fill_missing_dates(df, date_col='date', break_date=False)` (Static)

Fills gaps in dates using business day offsets and optionally extracts date components.

Parameters

- `df` (`pd.DataFrame`): Input with date column.
- `date_col` (`str`): Name of date column.
- `break_date` (`bool`): If True, adds `day`, `month`, `year` columns.

Returns

- `pd.DataFrame`: Filled DataFrame.

Algorithm

- Sorts by index.
- For NaN dates: Uses previous/next valid date + BDay offset.
- Fallback: Current timestamp if isolated.

Example

```
df['date'] = pd.to_datetime(df['date']) # Ensure datetime
df_filled = StockDataProcessor.fill_missing_dates(df, break_date=True)
print(df_filled[['date', 'day', 'month', 'year']].head())
```

`markov_impute(series, n_bins=20, strategy="mode")` (Static, Internal)

Markov chain-based imputation for series (used in `fill_nan_advanced`).

Parameters

- `series` (`pd.Series`): Input series.
- `n_bins` (`int`): Number of states.
- `strategy` (`str`): "mode" (deterministic) or "random" (stochastic).

Returns

- `pd.Series`: Imputed series.

Algorithm

- Bins non-NaN values.
- Builds transition matrix.
- Predicts next state from previous.

Example

```
imputed = StockDataProcessor.markov_impute(df['close'].dropna())
```

`fill_nan_advanced(df, col_tech_map)` (Static)

Applies sequential imputation techniques per column.

Parameters

- `df` (`pd.DataFrame`): Input DataFrame.
- `col_tech_map` (`dict`): {col: [(tech, params)]}. See method docstring for full list (e.g., "knn", "kalman").

Returns

- `pd.DataFrame`: Filled copy.

Supported Techniques

Technique	Description	Params Example
<code>drop</code>	Drop rows with NaNs in col	<code>{}</code>

Technique	Statistical fill	Params Example
ffill / bfill	Forward/backward fill	{}
sma / rolling	Simple moving average	{'window': 14}
ema	Exponential moving average	{'alpha': 0.3}
linear / quadratic / cubic	Polynomial interpolation	{}
knn	K-Nearest Neighbors	{'n_neighbors': 3}
markov	Markov chain	{}
weighted_combo	0.5SMA + 0.5EMA	{'window': 14, 'alpha': 0.3}
kalman	Kalman smoothing	{}

Raises

- ValueError: Unknown technique.

Example

```
col_tech_map = {
    "close": [{"knn", {"n_neighbors": 5}}, {"kalman", {} }],
    "volume": [{"sma", {"window": 10}}]
}
df_filled = StockDataProcessor.fill_nan_advanced(df, col_tech_map)
```

plot\_graph(df, columns, graph\_type, size=(10,6), color='blue', stacked=False) (Static)

Generates diverse plots for EDA.

Parameters

- df (pd.DataFrame): Input data.
- columns (list[str]): Columns to plot (varies by type).
- graph\_type (str): See table below.
- size (tuple): Figure size.
- color (str/list): Color scheme.
- stacked (bool): For area plots.

Supported Graph Types

Type	Columns Req.	Library	Use Case
line / scatter	2 (x,y numeric)	Matplotlib	Trends/Correlations
bar	1-2	Matplotlib	Counts/Categories
hist	1 (numeric)	Matplotlib	Distributions
box / violin	1 (numeric)	Seaborn	Outliers/Summary
pairplot	>=2	Seaborn	Multi-var EDA
area / stacked_area / stream	>=2	Matplotlib	Cumulative Trends
pie	1-2	Matplotlib	Proportions
waterfall	2	Matplotlib	Cumulative Changes
candlestick	4+ OHLC	mplfinance	Price Action

treemap / sunburst Type	2+ Columns Req.	Plotly Library	Hierarchies Use Case
choropleth	2 (region,value)	Plotly	Geo-Maps

Returns

- None (displays plot).

Raises

- ValueError : Invalid columns/type.

Example

```
StockDataProcessor.plot_graph(df, ["date", "close"], "line")
StockDataProcessor.plot_graph(df, ["open", "high", "low", "close"], "candlestick", size=(12,8))
```

`detect_outliers_advanced(df, numeric_cols, z_thresh=3, mod_z_thresh=3.5, rolling_window=5, price_change_thresh=0.05, plot_graphs=True, combine='union', vote_thresh=None)` **(Static)**

Detects outliers using 7 methods.

Parameters

- df (pd.DataFrame): Input.
- numeric\_cols (list[str]): Columns to analyze.
- Detection thresholds as named.
- combine (str): "union"/"intersection".
- vote\_thresh (int): Min methods to flag (overrides combine).
- plot\_graphs (bool): Generate plots (box, hist, scatter, violin).

Detection Methods

Method	Description	Threshold
z_score		Z
modified_z	Robust Z via MAD	Median ± MAD
iqr	Outside 1.5*IQR	Q1-Q3
rolling	Outside rolling mean ± 3*std	Window-based
price_change		% change
returns_z	Z on returns	Returns
cusum	Cumulative shifts	>3*std

Returns

- dict: {col: {per\_method': {method: set(indices)}, 'combined': set(indices)}}.

Example

```
outliers = StockDataProcessor.detect_outliers_advanced(df, ["close"], plot_graphs=True, combine="intersection")
print(f"Outliers in close: {len(outliers['close']['combined'])}")
```

`treat_outliers(df, outlier_results, treatment_map)` **(Static)**

Applies treatments to detected outliers.

Parameters

- `df` (`pd.DataFrame`): Input.
- `outlier_results` (dict): From `detect_outliers_advanced`.
- `treatment_map` (dict): {col: [(method, params)]}.

Supported Treatments

Method	Description	Params Example
<code>delete</code>	Drop outlier rows	<code>{}</code>
<code>winsorize</code> / <code>cap</code>	Clip to quantiles	<code>{'lower':0.01, 'upper':0.99}</code>
<code>median_replace</code>	Replace with median	<code>{}</code>
<code>mean_cap</code>	Clip to mean $\pm k \cdot \text{std}$	<code>{'k':3}</code>
<code>log_transform</code> / <code>sqrt_transform</code> / <code>boxcox</code>	Transformations	<code>{}</code> (clip negatives)
<code>robust_flag</code>	Add flag column	<code>{}</code>
<code>interpolate_linear</code> etc.	Interpolate	<code>{'method':'linear'}</code>
<code>rolling_mean</code> / <code>rolling_median</code>	Smoothing	<code>{'window':5}</code>
<code>ema_smooth</code>	EMA	<code>{'alpha':0.3}</code>
<code>kalman</code>	Kalman filter	<code>{}</code>
<code>markov_prev</code> / <code>markov_avg</code>	Markov replacements	<code>{}</code>

Returns

- `pd.DataFrame`: Treated copy.

Raises

- `ValueError`: Unknown method.

Example

```
treatment_map = {"close": [("winsorize", {"lower":0.05, "upper":0.95}), ("robust_flag", {})]}
df_treated = StockDataProcessor.treat_outliers(df, outliers, treatment_map)
```

`test_data()` (Static)

Generates sample OHLCV DataFrame (20 business days, Oct 2025).

Returns

- `pd.DataFrame`: Columns: `date`, `low`, `high`, `open`, `close`, `volume`.

Example

```
test_df = StockDataProcessor.test_data()
print(test_df.describe())
```

# Usage Examples: End-to-End Workflow

## 1. Basic Data Pipeline

```

# Load test data
df = StockDataProcessor.test_data()

# Introduce NaNs for demo
df.loc[5:7, "close"] = np.nan

# Fill dates (already complete in test)
df = StockDataProcessor.fill_missing_dates(df)

# Impute NaNs
col_tech_map = {"close": [{"linear", {}}]}
df = StockDataProcessor.fill_nan_advanced(df, col_tech_map)

# Visualize
StockDataProcessor.plot_graph(df, ["date", "close"], "line")

# Outliers
outliers = StockDataProcessor.detect_outliers_advanced(df, ["volume"], price_change_thresh=0.1)
print(outliers)

```

## 2. Real-World: AAPL Analysis with Outliers

```

df = StockDataProcessor.download_stock_data("AAPL", "1y")
df = StockDataProcessor.fill_missing_dates(df)

# Detect outliers in returns
df["returns"] = df["close"].pct_change()
outliers = StockDataProcessor.detect_outliers_advanced(
    df, ["returns"], z_thresh=2.5, plot_graphs=False, vote_thresh=3
)

# Treat: Smooth returns
treatment_map = {"returns": [{"ema_smooth", {"alpha": 0.2}}]}
df_treated = StockDataProcessor.treat_outliers(df, outliers, treatment_map)

# Plot treated data
StockDataProcessor.plot_graph(df_treated, ["date", "returns"], "scatter", color="red")

```

## 3. Advanced Imputation Chain

```

col_tech_map = {
    "open": [{"ffill", {}}, {"knn", {"n_neighbors": 4}}],
    "volume": [{"markov", {}}, {"kalman", {}}]
}
df_imputed = StockDataProcessor.fill_nan_advanced(df, col_tech_map)
StockDataProcessor.plot_graph(df_imputed, ["date", "volume"], "area", stacked=True)

```

## Performance Considerations

- **Scalability:** Handles up to 10k rows efficiently; for larger datasets, subsample or parallelize imputations.
- **Memory:** Copies DataFrames; use `inplace=True` where possible in custom extensions.
- **Compute-Intensive Methods:** KNN/Kalman scale  $O(n^2)/O(n)$ ; limit to key columns.

## Contributing

1. Fork the repo.
2. Create a feature branch ( `git checkout -b feature/amazing-feature` ).
3. Commit changes ( `git commit -m 'Add amazing feature'` ).
4. Push and open a PR.

Run tests: `pytest tests/` . Lint: `black .` .

## Acknowledgments

- Built with love by [QUANT-FINANCE COMMUNITY , NIT ROURKELA].
- Thanks to yfinance AND the open-source community.

For issues or features, open a GitHub issue. Let's make stock analysis accessible!